

Topic 1: A Better UDP

Carson Clanton

Department of Computer
Science and Engineering
The University of Texas
at Arlington
Arlington, Texas 76019
Email: clanton@uta.edu

Kevin Hayes

Department of Computer
Science and Engineering
The University of Texas
at Arlington
Arlington, Texas 76019
Email: @mavs.uta.edu

William Hunter

Department of Computer
Science and Engineering
The University of Texas
at Arlington
Arlington, Texas 76019
Email: @mavs.uta.edu

Richard Meth

Department of Computer
Science and Engineering
The University of Texas
at Arlington
Arlington, Texas 76019
Email: @mavs.uta.edu

Ryan Nordeen

Department of Computer
Science and Engineering
The University of Texas
at Arlington
Arlington, Texas 76019
Email: @mavs.uta.edu

Abstract—This article describes an approach to improving the reliability of UDP. Our approach uses sequence numbers for packets to force in-order delivery to the Network Application layer and Error-Correcting Codes to help reproduce lost packets without requesting retransmission of these missing packets. The level of reproducibility will be largely dependent on the Error-Correcting Code chosen and the selection of its associated redundancy parameters.

I. INTRODUCTION

When a software engineer begins designing a network application, he or she has to make a decision between two transport layer options, TCP or UDP. There are advantages and disadvantages to both. With UDP comes nimbleness and low overhead, along with dropped and possibly out-of-sequence packets. With TCP comes reliable stream-based communication at the expense of throughput. There isn't much middle ground between the two. The approach discussed in this paper attempts to change that.

We don't attempt or intend to redesign either protocol. Instead, we seek to enhance the reliability of UDP, so that it is a little more like TCP, while remaining connectionless. To do this, we will require two added features to be built "on top" of UDP: Sequence Numbers and Error-Correcting Codes.

The introduction of sequence numbers should be somewhat obvious; on the receiving end of the UDP pathway there should be some way to put packets back in order. The subtlety, however, of requiring sequence numbers is that in order for our error-correcting codes approach to work, for the code we experimented with, at least 75% of the packet payload data must be available and properly sequenced. So, the inclusion of sequence numbers is twofold.

The second feature to be added, Error-Correcting Codes or ECC, is a little more complicated and requires more of an introduction, if for no other reason than to get acclimated to the terminology in the literature.

Given the no-guaranty and low-reliability characteristics of the networking layer and the little added value of UDP, a virtual communication link through a UDP socket represents an unreliable, and potentially noisy, channel, making it a good candidate for a ECC solution.

A. Error-Correcting Codes

Below in Figure 1 is a schematic of the general communication link for sending messages. Here the message




Fig. 1. Schematic of a general communication link

is first sent by the *source* to the *encoder* where the message is assigned a *codeword*, i.e. a string of characters from some chosen alphabet. The encoded message is sent through a *channel*, which is certain to have some level of noise that may possibly alter the codeword. This possibly-altered codeword is then received by the *decoder* where it is matched with its most likely candidate codeword which is in turn translated into a message resembling the original message and passed on to the *receiver*. The degree of resemblance to the original message depends on how appropriate the code is in relation to the channel.

Example I.1. Suppose we are using an alphabet containing only two symbols, say 0 and 1 and we wish to send a message to a friend of either "YES" or "NO". We first might consider encoding 1 for "YES" and 0 for "NO"; however, if the codeword 0 is sent and then altered by the channel to be received as 1, the decoder will incorrectly pass the message "YES" to the receiver. One way of improving the situation is to add redundancy by repeating the symbol. For instance, we might encode "YES" as the codeword 11 and "NO" as 00. Now if "NO" is sent, two errors would have to occur before the decoder returns an incorrect message. If one error occurs then the received vector will either be 10 or 01, neither of which is a codeword. At this point, the receiver might request a retransmission. This is an example of a code in which one error may be detected. To strengthen the code to be error-correcting we might increase the redundancy by repeating the message five times. Thus we encode our message "NO" as 00000. Perhaps the channels interferes to cause the decoder to receive the vector 00110. Using the method of

nearest neighbor decoding the decoder assesses the message and decides that of the two possible codewords (i.e. 00000 and 11111) 00000 is more likely the one originally sent and hence is correctly decoded as “NO”.

Definition I.2. Let \mathbf{F} be a finite set, or **alphabet**, of q elements. A q -ary code \mathbf{C} is a set of finite sequences of symbols of \mathbf{F} , called **codewords** and written $x_1x_2\cdots x_n$, or (x_1, x_2, \dots, x_n) , where $x_i \in \mathbf{F}$ for $i = 1, \dots, n$. If all the sequences have the same length n , then \mathbf{C} is a **block code** of **block length** n .

Given an alphabet \mathbf{F} , it is consistent with terminology for vector spaces when \mathbf{F} is a field to denote the set of all sequences of length n of elements of \mathbf{F} by \mathbf{F}^n and to call these sequences vectors. The member of \mathbf{F} at the i th position of a vector is known as the coordinate at i .

Definition I.3. Let $v = (v_1, v_2, \dots, v_n)$ and $w = (w_1, w_2, \dots, w_n)$ be two vectors in \mathbf{F}^n . The **Hamming distance**, $d(v, w)$, between v and w is the number of coordinate places in which they differ:

$$d(v, w) = \{i | v_i \neq w_i\}$$

We will usually refer to the Hamming distance as simply the **distance** between two vectors.

Nearest-neighbor decoding is a method in which the received vector is translated as the codeword of smallest distance, whenever it is uniquely determined. This method maximizes the likelihood of correcting errors provide the two assumptions mentioned above hold (i.e. probability of an error $< 1/2$ and each symbol is equally likely to be transmitted).

Definition I.4. The **minimum distance** $d(\mathbf{C})$ of a code \mathbf{C} is the smallest of the distances between distinct codewords; i.e.

$$d(\mathbf{C}) = \min\{d(v, w) | v, w \in \mathbf{C}, v \neq w\}$$

Example I.5. The error-correcting code in Example I.1 is a binary code of block length 5. The distance between the codeword 00000 and the received vector 00110 is 2. The distance between the same received vector and the codeword 11111 is 3. Thus, 00000 is the nearest neighboring codeword. The minimum distance of this code is 5.

Definition I.6. A code \mathbf{C} over \mathbf{F} of length n is **linear** if \mathbf{C} is a subspace of the vector space $V = \mathbf{F}^n$. If $\dim(\mathbf{C}) = k$ and $d(\mathbf{C}) = d$, then we write $[n, k, d]$ for the q -ary code \mathbf{C} ; if the minimum distance is not specified we simply write $[n, k]$. The **information rate** is k/n and the **redundancy** is $n - k$.

Since \mathbf{C} is a subspace, it follows that the **all-zero** vector must be a codeword in \mathbf{C} also the sum of any two codewords must also be a codeword in \mathbf{C} . We will also see after the next definition that the distance d is also easy to calculate for linear codes.

Definition I.7. Let $V = \mathbf{F}^n$. For any vector $v = (v_1, v_2, \dots, v_n) \in V$ set

$$S = \{i | v_i \neq 0\}$$

Then S is called the **support** of v and the **weight** of v is $|S|$. The **minimum weight** of a code is the minimum of the weights of the nonzero codewords.

Example I.8. Listed below are all 16 codewords in the linear $[7, 4, 3]$ binary code. They are listed in order of increasing weight.

0000000	0111000	1000111	0101101
1100100	0001110	1101010	0110110
1010010	0010101	1110001	0011011
1001001	0100011	1011100	1111111

Notice that the minimum weight is 3. By Theorem 1.1 it follows that $d(\mathbf{C}) = 3$. It can be shown that these 16 codewords form a subspace of \mathbf{F}^7 with basis $\{(1, 0, 0, 0, 1, 1, 1), (0, 1, 0, 0, 0, 1, 1), (0, 0, 1, 0, 1, 0, 1), (0, 0, 0, 1, 1, 1, 0)\}$. Thus $\dim(\mathbf{C}) = 4$ as conveyed in the notation and the redundancy is 3 since $n - k = 7 - 4 = 3$.

Theorem I.9. Let \mathbf{C} be a code of minimum distance d . If $d \geq 2t + 1$, then \mathbf{C} can be used to correct up to t errors.

II. DESIGN

A Better UDP is “wedged in” as another IP stack layer sitting between the transport layer and the application layer, see Figure 2. We did not modify any kernel source and A Better UDP is not a standard yet, so to use it an application developer or software engineer will have to include this functionality as a library. Instead of using a socket library to interface with the transport layer, the engineer will link with A Better UDP library to interface with the protocol stack below the application layer. Henceforth we will refer to A Better UDP as BUDP.




Fig. 2. Five-layer Internet Protocol Stack with a Better UDP Wedged in

Messages, from the application layer, are passed to BUDP, where, currently, four at a time are buffered, assigned sequence numbers, encoded to produced an ECC, and then sent as datagrams via standard UDP to their destination address. At the destination, the datagrams are collected and stored in an accumulation buffer, where they are placed in their proper sequence, see Figure 3. One of three conditions, prompts their delivery to the application layer at the destination:

- 1) All four datagrams are received intact.
- 2) Three datagrams are received as well as the ECC, intact.
- 3) A countdown timer reaches zero.

The first of these conditions is obvious, while the other two require further explanation. It is easier to explain 3) first, so we will do that now. As stated previously in our introduction, we don't intend to rewrite TCP and we would like to keep our UDP implementation connectionless, so instead of requesting retransmission of missing datagrams, for example, when the one of the first two conditions above are not met, our implementation will simply give what it has to the application layer at the destination. Our justification for doing this is that it is what a standard UDP implementation would do. Actually, a standard implementation of UDP would do less, because our implementation will at least give the application layer in-order messages.

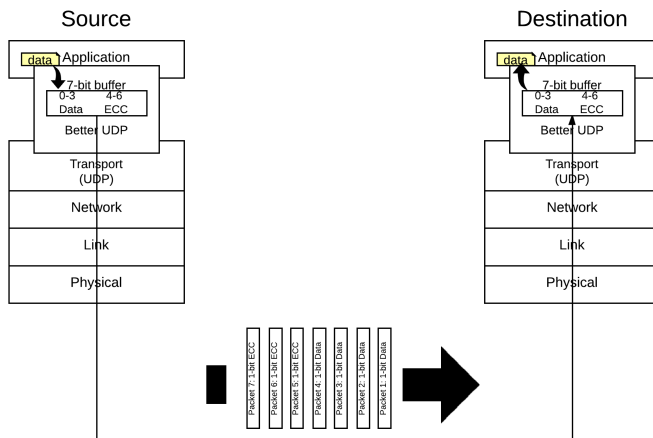


Fig. 3. A Better UDP Overview Schematic

The ECC is also sent in follow-on packets to the destination. They also contain sequence numbers, which are a continuation of the sequence numbers assigned to data packets. The purpose of buffering is to support the process of encoding the ECC.

For the purposes of prototyping our design of a better UDP, we will use the single error-correcting Hamming code introduced above. It is quite limited in what it can be used for, since it only corrects one error, because for our purposes it will do just fine.

III. CONCLUSION

The conclusion goes here.

ACKNOWLEDGMENT

The authors would like to thank...

REFERENCES

- [1] H. Kopka and P. W. Daly, *A Guide to L^AT_EX*, 3rd ed. Harlow, England: Addison-Wesley, 1999.